## Exercise 1: SVD and its properties

**Data:**

Step 1: Generate two random matrices: A = randn(1000,2); B = randn(200,2);

Step 2: Construct matrix X = A*B’ + noise, i.e., where noise is =0.1*randn(1000,200)

- What is the rank of X?
- Plot the singular values of X
- Is X low-rank?
- What is the best rank-2 approximation of X?
- Compute what percent of the “energy” is explained using the best rank-2 approx.
- Plot the significant left and right singular vectors.

# Exercise 2: Recommender Systems

Given the following user-movie rating matrix, can you group users based on their interests in movies? Which groups of users are interested in what type of movies?

| | Aloha | Star Wars | American Pie | Hunger Games | Silver Linings | Maze Runner |
|---|---|---|---|---|---|---|
| User 1 | 5 | 2 | 4 | 1 | 3 | 2 |
| User 2 | 1 | 4 | 1 | 4 | 1 | 4 |
| User 3 | 3 | 1 | 5 | 2 | 5 | 2 |
| User 4 | 1 | 4 | 1 | 5 | 3 | 4 |
| User 5 | 3 | 1 | 4 | 2 | 3 | 2 |
| User 6 | 2 | 4 | 2 | 4 | 2 | 4 |

**load exercise2_usermovie.mat**
**Compute its SVD**

# Exercise 3: Text Mining

Given the following document-term matrix, find the related documents and what they are about.

**Document 1:** **Tiger stopped playing golf**

**Document 2:** **News** about **Tiger** and his **golf career**

**Document 3:** **Golf career** of **Tiger** in **jeopardy**

**Document 4:** **Tiger** and his **wife** in the **news**

**Document 5:** The **new zoo featuring** the **big cat family**: **tigers** and **lions**

**Document 6:** **Tigers** – the **big cats** – in the **new zoo**

**Document 7:** **Tigers** and **lions**, **which** are the **biggest cats**?

| | tiger | stop | play | golf | news | career | jeopardy | wife | zoo | featuring | new | big | cat | family | lions | which |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| Doc1 | 1 | 1 | 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| Doc2 | 1 | 0 | 0 | 1 | 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| Doc3 | 1 | 0 | 0 | 1 | 0 | 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| Doc4 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| Doc5 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 0 |
| Doc6 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 1 | 1 | 0 | 0 | 0 |
| Doc7 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 1 | 1 |

**load exercise3_docterm.mat**
**Compute its SVD**

Which group does this document belong to?

**Document 8:** **Tiger** back to **golf**

$$\mathbf{q}^T = \begin{bmatrix} 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

# Exercise 4: Metabolomics

Some are fed with 10g apple while some are controls (no apple)

Liquid Chromatography-
Mass Spectrometry (LC-MS)

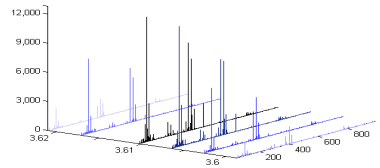


**features**

**samples**


**load exercise4_metabolomics.mat**

- Preprocess the data
- Plot singular values
- Compute the best rank-r approximation that represents the 80% of the “energy”
- Can you group the samples based on apple consumption (Class1:10g apple, Class2: no apple)?

simula

# Exercise 5: Link Prediction using the DBLP data

This exercise is about using the Singular Value Decomposition for the temporal link prediction problem. When you load *hw1_data.mat*, you will see two data sets, i.e., $\mathcal{X}$ and $\mathbf{Y}$. $\mathcal{X}$ shows the number of papers published by authors at various conferences between 1991 and 2004. It is of size 471 (authors) $\times$ 366 (conferences) $\times$ 14 (years). Given $\mathcal{X}$, we want to predict who is going to publish at which conference in 2005. Matrix $\mathbf{Y}$ shows the ground truth, i.e., publications in 2005.

Execute the following steps and return the outputs in the deliverables:

- Change each nonzero entry of $\mathcal{X}$ as $x_{ijk} = log(x_{ijk}) + 1$, where $x_{ijk} \neq 0$.
- Collapse the three-way array $\mathcal{X}$ by summing up over the years mode and form an authors by conferences matrix of size 471 $\times$ 366. Let this matrix be $\mathbf{Z}$.
- Compute the SVD of $\mathbf{Z}$.
- Construct the best rank-$K$ approximation of $\mathbf{Z}$ denoted as $\hat{\mathbf{Z}}_K$ for different $K$ values, i.e., $K = \{2, 10, 20, 50, 100, 300\}$. Entries of $\hat{\mathbf{Z}}_K$, i.e., $\hat{\mathbf{Z}}_K(i, j)$ can be used as scores to predict if there is a link between the $i^{th}$ author and $j^{th}$ conference in 2005. A link means an author publishes at a conference.
- Replace every nonzero entry of $\mathbf{Y}$ with 1. Vectorize $\mathbf{Y}$, i.e., $\mathbf{Y}(:)$ in MATLAB notation, which will correspond to the true labels (0's and 1's).
- For each value of $K$, vectorize $\hat{\mathbf{Z}}_K$, which will correspond to the scores/predictions.
- For each value of $K$, plot the Receiver Operating Characteristics (ROC) Curve and calculate the area under the curve (AUC) (Note: You can use the perfcurve function in MATLAB).

simula

# Exercise 6: Solving Least Squares using SVD

simula